

BART-Survival: A Bayesian machine learning approach to survival analyses in Python

Jacob Tiegs^{1,2}, Julia Raykin¹, and Ilia Rochlin¹

¹ Inform and Disseminate Division, Office of Public Health Data, Surveillance, and Technology, Centers for Disease Control and Prevention, Atlanta, GA, USA ² Metas Solutions, Atlanta, Georgia
Corresponding author

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#)
- [Repository](#)
- [Archive](#)

Editor: [Open Journals](#)

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

BART-Survival is a Python package that allows time-to-event (survival) analyses in discrete-time using the non-parametric machine learning algorithm, Bayesian Additive Regression Trees (BART). BART-Survival combines the performance of the BART algorithm with the complementary data and model structural formatting required to complete the survival analyses. The library contains a convenient application programming interface (API) that allows a simple approach when using the library for survival analyses, while maintaining capabilities for added complexity when desired. The package is intended for analysts exploring use of flexible non-parametric alternatives to traditional (semi-)parametric survival analyses.

Statement of need

Survival analyses are a cornerstone of public health and clinical research in such diverse fields as cancer, cardiovascular disease, and infectious diseases (Altman & Bland, 1998; Bradburn et al., 2003). Traditional parametric and semi-parametric statistical methods, such as the Cox proportional hazards model, are commonly employed for survival analyses (Cox, 1972). However, these methods have several limitations, particularly when applied to complex data. One major issue is the need for restrictive assumptions, such as proportional hazards and predefined functional forms, which may not hold true in complex, real-world healthcare data (Harrell, 2015; Ishwaran et al., 2008). Additionally, these methods often struggle with high-dimensional datasets, leading to problems with overfitting, multicollinearity, and dealing with complex interactions (Ishwaran et al., 2008; Joffe et al., 2013).

More recently, non-parametric machine learning approaches have been introduced to address these limitations by reducing the need for restrictive assumptions and providing increased capabilities for more accurately modeling underlying distributions and complex interactions (Harrell, 2015; Ishwaran et al., 2008). BART is one such machine learning method that has demonstrated utility in the survival setting through its performance in identifying underlying statistical distributions (Chipman et al., 2010; R. Sparapani et al., 2021). BART offers flexibility in modeling complex relationships and interactions within the data without requiring the specification of a particular functional form (R. A. Sparapani et al., 2016).

Currently, the only BART survival algorithm readily available exists as part of the BART R package, which contains a library of various BART-based approaches in addition to a BART survival analysis application (R. Sparapani et al., 2021; R. A. Sparapani et al., 2016). BART-Survival package described here combines the survival analysis approach outlined in the BART R package with the foundational Python-based probabilistic programming language library, PyMC (Abril-Pla et al., 2023), and the accompanying BART algorithm from the PyMC-BART library (Quiroga et al., 2023). Our aim in developing BART-Survival is to provide accessibility

to the BART survival algorithm within the Python programming language. This contribution is beneficial for analysts when Python is the preferred programming language, the analytic workflow is Python-based, or when the R language is unavailable for analyses. Additionally, BART-Survival package abstracts away the complexities of the PyMC and PyMC-BART libraries through use of a pre-specified core model and generalized functionality that can accommodate analyses across various survival settings. The BART-Survival package is intended for public health and clinical professionals and students who are looking for non-parametric alternatives to traditional (semi-)parametric survival analysis, especially for use in large, complex healthcare data and machine learning applications.

Methods

The following sections provides details on the methods employed by the BART-Survival library, focusing specifically on the discrete-time Survival algorithm used. For review of the BART algorithm we refer to associated PyMC-BART publication (Quiroga et al., 2023).

Background

The BART-Survival package provides a discrete-time Survival method which aims to model Survival as a function of a series of probabilities that indicate the probability of event occurrence at each discrete time. For clarity, the probability of event occurrence at each discrete time will be referred to as the risk probability.

In combination with a structural configuration of the data, the discrete-time algorithm allows for flexible modeling of the risk probabilities as a non-parametric function of time and observed covariates. The series of probability risks can then be used to derive the Survival probabilities, along with other estimates of interest.

The foundation of the method is simple.

1. Starting with a simple event-time dataset, create a sequence of time intervals that represent the unique discrete-time intervals observed in the data. Each interval is represented as a t_j , where $j = 1, \dots, k$ and k is the length of the set of unique observed times.

For example if the observed event-time data is:

event status	event time
1	1
0	2
1	2
1	4
1	4
1	5

Then the set of unique observed times is $[1, 2, 4, 5]$ with $k = 4$ indices and the corresponding t_j intervals are:

index	interval
t_1	$[0, 1)$
t_2	$[1, 2)$
t_3	$[2, 4)$
t_4	$[4, 5)$

Where the intervals are denoted with the properties, “[” indicating exclusive times and “)” indicating inclusive times in an interval. Additionally, the **event status** column is defined as (1 = event, 0 = censor)

2. Then with the constructed time intervals t_1, \dots, t_k , tally the following for each interval:

- the number of observations with an event
- the number of observation censored
- the total number of observations eligible to have an event (at risk) at the start of the interval

Continuing the above example the corresponding frequencies for each interval are:

index	interval	event	censor	at risk
t_1	$[0, 1)$	1		6
t_2	$[1, 2)$	1	1	5
t_3	$[2, 4)$	2		3
t_4	$[4, 5)$	1		1

3. Finally, the risk of event occurrence within each interval t_j can simply be derived as:

$$P_{t_j} = \frac{n \text{ events}_{t_j}}{n \text{ at risk}_{t_j}}$$

and the Survival probability at a time-index j , can be derived as:

$$S(t_j) = \prod_{q=1}^j (1 - P_{t_q})$$

Applied to our example the risks of event P_{t_j} are:

index	interval	event	censor	at risk	P_{t_j}
t_1	$[0, 1)$	1		6	0.167
t_2	$[1, 2)$	1	1	5	0.2
t_3	$[2, 4)$	2		3	0.667
t_4	$[4, 5)$	1		1	1.0

And the corresponding Survival estimates at times $[1, 2, 4, 5]$ are:

index	time	Survival
t_1	1	.83
t_2	2	.67
t_3	4	.22
t_4	5	0

BART-Survival builds off this simple foundation by replacing each naively derived P_{t_j} with a probability risk prediction, $p_{t_j|x_{ij}}$ yielded from the BART regression model. The x_{ij} values are the associated covariate values for each observation i at each discrete time j . The predicted values $p_{t_j|x_{ij}}$ can be further transformed into Survival probability estimates or other estimands of interest. Formally, the Survival probability estimate for a single observation i at time-index j can be derived as:

$$S(t_j|x_i) = \prod_{q=1}^j (1 - p_{t_q|x_{iq}})$$

91 Data Preparation

92 As displayed in the above example, Survival data is typically given as paired (*event status*,
93 *event time*) outcomes, along with a set of covariates x_i for each observation. In this setup,
94 *event status* is typically a binary variable, with 1 = **event** and 0 = **censored**.

95 Obtaining the $p_{t_j|x_{ij}}$ values and subsequent Survival estimates is a two-step process. Each step
96 of the process requires generating a specialized augmented datasets from the generic Survival
97 data. The first step in the process involves training the BART model which requires the
98 **training augmented dataset (TAD)**. The second step involves generating the $p_{t_j|x_{ij}}$ predictions
99 from the trained BART model, which requires the **predictive augmented dataset (PAD)**.

100 The **TAD** is created by transforming each single observation row into a series of rows representing
101 the observation over a time series. When transformed a single observation's information is
102 represented as a series of tuples, each tuple containing the augmented values of *status* and
103 *time*, as well as replicates of the covariates x . The series of tuples is constructed using the
104 k distinct times from the generic dataset. These times are represented by t , which can be
105 indexed by j , where j is in the set $1, \dots, k$. Then for each observation's *event time* _{i} there is an
106 index j in $1, \dots, k$ where $t_j = \text{event time}_i$. This special index j is denoted as m_i and can be
107 defined as:

$$t_{m_i} = t_j = \text{event time}_i$$

108 For each observation, a set of tuples is created containing a total of m_i tuples which are used
109 to represent that observation's information. Each tuple is created for a specific time t_j where
110 $j = 1, \dots, m_i$. Additionally, for all tuples with times $t_j < t_{m_i}$ the value of *status* is set to 0
111 and for the time $t_j = t_{m_i}$ the value of *status* is set to the value *event status* _{i} . Covariate
112 information is simply replicated across tuples.

113 For example if the Survival data is:

observation	status	time	with	unique times (k=6)
1	1	4		— — —
2	0	6		4
3	1	6		6
4	1	7		7
5	1	8		8
6	0	8		12
7	1	12		14
8	1	14		

114 Then the transformation for observation 7 would be represented in the *augmented* dataset as
115 the sequence of rows:

observation	status	time
7	0	4
7	0	6
7	0	7
7	0	8
7	1	12

116 Similarly the transformation for observation 6 would be represented as:

observation	status	time
6	0	4
6	0	6
6	0	7
6	0	8

117 It is important to reiterate that for each observation i , the new set of rows created only include
 118 the time points $t_j \leq t_{m_i}$. For observation 7 the $t_{m_i} = 12$ and the corresponding times in
 119 the **TAD** are 4, 6, 7, 8, 12. For observation 6 the $t_{m_i} = 8$ and the corresponding times in the
 120 **TAD** are 4, 6, 7, 8.

121 The utility of the **TAD** is that it unlinks *event time* from the *event status*. In this setup, the
 122 constructed *status* variable (which is a simple binary variable taking values 1 or 0) represents
 123 the outcome and *time* is a covariate. Then treating each row of the **TAD** as an independent
 124 observation, the outcome can be modeled as a probit regression of *status* over *time* and any
 125 additional covariates x . The trained regression model from this dataset can be used to yield
 126 probability predictions for each time t_j for j in $1, \dots, k$. These probability predictions hold the
 127 interpretation as being the probability risk of event occurrence at each time conditional on the
 128 event not having already occurred at a previous time. The predictions for each observation i
 129 and time t_j can be used to yield the various analytic targets, such as Survival probabilities.

130 To generate the predictions of probability risks over the observed times the **PAD** dataset is used.
 131 The **PAD** transformation is similar to the **TAD**, but is simpler in that for each observation a
 132 set of size k tuples is created. Each tuple in the set contains the time t_j for $j = 1, \dots, k$ and
 133 any additional covariates x .

134 Continuing the examples above, there are $k = 6$ the unique times. For $j = 1, \dots, k$ the resulting
 135 distinct times are $t_j = 4, 6, 7, 8, 12, 14$ and the **PADs** for observations 6, and 7 would be:

observation	time
6	4
6	6
6	7
6	8
6	12
6	14
7	4
7	6
7	7
7	8
7	12
7	14

136 Notably the **TAD** and **PAD** lengths can differ. **PAD** length is simple to calculate. If n is
 137 the number of observations in the generic dataset and k is the number of the unique times
 138 observed in the generic dataset, then the length of the **PAD** is simply $n * k$. In the example
 139 data above it would be $(8 * 6) = 36$.

140 The **TAD** length can be calculated using each observation's m_i , which again is defined as
 141 being the index that resolves $t_{m_i} = \text{event time}_i$. The **TAD** length can be calculated as:

$$\text{TAD}_{\text{length}} = \sum_{i=1}^n m_i$$

142 This can be demonstrated with the example dataset where the **TAD** length can be found as
143 the sum of the m_i column which resolves to a **TAD** length of 27:

observation	event time _{<i>i</i>}	<i>m_i</i>
1	4	1
2	6	2
3	6	2
4	7	3
5	8	4
6	8	4
7	12	5
8	14	6
		27

144 While the **PAD** and **TAD** lengths are not important for generating the models, they are helpful
145 to keep in mind when completing an analysis. Specifically, the fact that the **TAD** and **PAD**
146 can be far larger in length than the generic dataset. For example the generic dataset used
147 above is only 8 rows in length, but the **TAD** and **PAD** dataset are 27 and 36 rows respectively.

148 When the length of the two datasets become large enough to make computation difficult it is
149 recommended to downscale the *event time* values, which will reduce the length of the two
150 augmented datasets. The downscaling algorithm we provide in the library takes the *event time*,
151 divides by a scaling factor and then takes the “ceiling” truncation to create the new *scaled*
152 *event time*. Below the downscaling algorithm is demonstrated with the example dataset using
153 a scaling factor of 4:

time	downscaled time	rescaled time
4	1	4
6	2	8
6	2	8
7	2	8
8	2	8
8	2	8
12	3	12
14	4	16

154 Downscaling the *event time* values causes loss of granularity in the set of discrete times
155 analyzed. As shown in the example above, after downscaling (and then rescaling) the analytic
156 targets can only be derived for the times 4, 8, 12 and 16, which is reduced from the original
157 set of times, 4, 6, 7, 8, 12, 14. While the analytic targets cannot be returned for the times
158 6, 7, 14, the training information contributed by the observations with these *event times* is
159 maintained through subsequent contribution at the downscaled times. Since there is no loss of
160 event information, the probability of event at the downscaled times 4, 8, 12, 16 will be equal to
161 the of the probability of event from a non-downscaled dataset at the select time 4, 8, 12, 16.

162 Downscaling can be safely applied to datasets without loss of precision of estimates at the
163 downscaled times. The only consideration required when using a downscaling procedure is to
164 ensure that the downscaled times granularity fulfills the needs of the analysis. For example, in a
165 multi-year health-outcomes study, downscaling from days to months could significantly reduce

computational burden without a meaningful loss of information in the outcome. However, if that same study reduced time from days to years, this could lead to loss of meaningful information in the outcome. The considerations for downscaling need to be made on a study-to-study basis and no general recommendation on downscaling factor can be provided.

Model

The BART-Survival algorithm is two-step algorithm. First, the **TAD** is used to train a non-parametric regression model of *status* on *time* and covariates. Then the **PAD** is used to yield predicted probabilities at each discrete time. These probability of event predictions can be used generate all additional estimates useful for statistical inference.

To motivate the use of the **TAD** in training the regression model an example **TAD** is displayed below. Here *status* can be represented as the outcome vector of y_{ij} values and *time* is represented by the t_{ij} values. Additional covariates are represented by a vector of x_{ij} values. The subscript i refers to the i^{th} observation of the generic dataset and the subscript j refers to the j^{th} index of the unique set of discrete times previously defined as t . In the **TAD** and **PAD**, the i, j indices can be thought of as multi-indices over the rows.

i	j	y	t	x
1	1	0	3	x_{11}
1	2	0	4	x_{12}
1	3	1	6	x_{13}
2	1	0	3	x_{21}
2	2	1	4	x_{22}
3	1	0	3	x_{31}
3	2	0	4	x_{32}
3	3	1	6	x_{33}

When using the **TAD**, each y_{ij} value can then be treated as independent draw of a *Bernoulli* distribution parameterized with the probability of event p_{ij} for observation i at time index j . The p_{ij} values are collected as the latent values yielded from a *probit* regression of y on t and x . Formally the model is defined as:

$$\begin{aligned}
 y_{ij} | p_{ij} &\sim \text{Bernoulli}(p_{ij}) \\
 p_{ij} | \mu_{ij} &= \Phi(\mu_{ij}) \\
 \mu_{ij} &\sim \text{BART}(t_{ij}, x_{ij})
 \end{aligned}$$

Where Φ is the normal cumulative distribution function and BART is the ensemble of regression trees which yield the μ_{ij} value for a given t_{ij}, x_{ij} combination.

Regarding the BART algorithm, in brief BART is a Bayesian approach to the ensemble regression trees class of non-parametric models. BART uses a combination of Bayesian priors placed on the tree generating components and a Markov Chain Monte Carlo sampling algorithm to iteratively generate tree ensembles which are collected as samples of the posterior distribution. The samples of the posterior distribution are subsequently used to generate distinct posterior predictive distributions for a given set of observations.

This implies that for each t_{ij}, x_{ij} combination, the algorithm first returns a distribution of p_{ij} values, denoted as $p_{ij_{dist}}$. Each value in the $p_{ij_{dist}}$ is a prediction mapped from a single tree ensemble of the posterior distribution. Point estimates and uncertainty intervals can be easily obtained from the $p_{ij_{dist}}$ as simple estimates from that distribution. For example the mean can be derived as the empirical average of $p_{ij_{dist}}$. Similarly, the even-tailed 95% credible

interval can be derived as the 5th and 95th percentiles of the $p_{ij_{dist}}$. Of note any reference to p_{ij} used in this article is referring to the mean point estimate of the $p_{ij_{dist}}$.

Further details regarding the BART implementation used in the BART-Survival algorithm can be found in PyMC-BART repository and the accompanying publication (Quiroga et al., 2023)

Survival

A trained BART-Survival model is used along with the **PAD** to generate the vector of p_{ij} predictions. To generate Survival estimates the p_{ij} vector is grouped by the i indices. Then using the following equation the Survival probability at discrete times can be constructed:

$$S(t_q|x_{iq}) = \prod_{j=1}^q (1 - p_{ij})$$

Where q is the index of the time for the desired estimate. As described in the previous section a full posterior predictive distribution is returned for each $S_{t_q, x_{iq}}$. Point estimates and credible intervals can be obtained through the empirical mean and percentile functions of the returned posterior predictive distribution for each $S(t_q|x_{iq})$.

Marginal Effects

BART-Survival provides capabilities to evaluate covariate effects through use of partial dependence functions which yield marginal effect estimates. The partial dependence function method involves generating predictions of p for the observations in a **partial dependence augmented dataset (PDAD)**. Using variations of the **PDAD** yields different sets of predictions p . These sets can then be contrasted to yield marginal effect estimates, including marginal Hazard Ratios which have similar interpretation as the Cox Proportional Hazard Model's conditional Hazard Ratios.

PDADs can be created through further augmentation of the previously described **PAD**. As a reminder, the **PAD** contains the generated time covariate t and k replicates of each x_i from the generic dataset, where k is the length of the uniquely observed *event times*. Starting with the **PAD**, the **PDAD** is generated through selection of a specific variable $x_{[I]}$ from the covariates x , and then deterministically setting $x_{[I]}$ to a specific value for all observations. The unselected covariates $x_{[0]}$ that are not augmented and are consistent with the values in the generic dataset. An example of creating a **PDAD** from the **PAD** is shown below. In this example a baseline **PAD** is used to create two **PDAD** datasets. The selected covariate x_2 is deterministically set to the values 0 or 1 for all observations in each dataset.

i	j	t	x_1	x_2	x_3
—	—	—	—	—	—
1	1	2	1.2	0	10
1	2	3	1.2	0	10
1	3	5	1.2	0	10
PAD: 2	1	2	2.4	1	12
2	2	3	2.4	1	12
2	3	5	2.4	1	12
3	1	2	1.9	0	3
3	2	3	1.9	0	3
3	3	5	1.9	0	3

i	j	t	x_1	x_2	x_3
1	1	2	1.2	1	10
1	2	3	1.2	1	10
1	3	5	1.2	1	10
2	1	2	2.4	1	12
2	2	3	2.4	1	12
2	3	5	2.4	1	12
3	1	2	1.9	1	3
3	2	3	1.9	1	3
3	3	5	1.9	1	3

227

i	j	t	x_1	x_2	x_3
1	1	2	1.2	0	10
1	2	3	1.2	0	10
1	3	5	1.2	0	10
2	1	2	2.4	0	12
2	2	3	2.4	0	12
2	3	5	2.4	0	12
3	1	2	1.9	0	3
3	2	3	1.9	0	3
3	3	5	1.9	0	3

228 From the two PDADs the predicted probabilities $p_{x_2=1}$, $p_{x_2=0}$ and Survival probabilities $S_{x_2=1}$,
229 $S_{x_2=0}$ can be generated:

i	j	$p_{x_2=1}$	$S_{x_2=1}$	$p_{x_2=0}$	$S_{x_2=0}$
1	1	.20	.80	.10	.90
1	2	.25	.60	.15	.77
1	3	.18	.49	.08	.70
2	1	.10	.90	.08	.92
2	2	.13	.78	.11	.81
2	3	.12	.69	.10	.74
3	1	.20	.80	.15	.92
3	2	.28	.58	.23	.82
3	3	.23	.44	.18	.74

230 The marginal expectations of $p_{x_{[i]}}$ and $S_{x_{[i]}}$ at a specific time can be further derived by taking
231 the average of the estimates over observations i for the specified time t indexed by j :

$$E_i[p_{x_{[i]}}|t_j] = \frac{1}{n} \sum_{i=1}^n p_{x_{[i]}}(t_j)$$

$$E_i[S_{x_{[i]}}(t_j)] = \frac{1}{n} \sum_{i=1}^n S_{x_{[i]}}(t_j)$$

233 Where E_i is the expectation over i, \dots, n observations. From the above example this yields the
234 expectations for time indices $j = 1, 2, 3$.

j	t	$E_i[p_{x_2=1}]$	$E_i[S_{x_2=1}]$	$E_i[p_{x_2=0}]$	$E_i[S_{x_2=0}]$
1	2	0.17	0.83	0.11	0.91
2	3	0.22	0.65	0.16	0.8
3	5	0.18	0.54	0.12	0.73

These expectations can be further used to make comparisons can be made between the evaluation of various values of $x_{[I]}$ across multiple **PDADs**. Common marginal effect estimates derived from these predicted values include:

- Marginal difference is Survival probability at time t :

$$\text{Surv. Diff.}_{\text{marg}}(t_j) = E_i[S_{x_{[I]_2}}(t_j)] - E_i[S_{x_{[I]_1}}(t_j)]$$

- Marginal Risk Ratio at time t :

$$\text{RR}_{\text{marg}}(t_j) = \frac{E_i[p_{x_{[I]_2}j}]}{E_i[p_{x_{[I]_1}j}]}$$

- Marginal Hazard Ratio (expectation over i and times t):

$$\text{HR}_{\text{marg}}(t_j) = \frac{E_{it}[p_{x_{[I]_2}j}]}{E_{it}[p_{x_{[I]_1}j}]}$$

Continuing the example, the marginal effect of x_2 as measured by difference in Survival probability when $x_2 = 1$ and $x_2 = 0$ at the times 2, 3, 5 can be examined below:

j	t	Surv.Diff.
1	2	−0.08
2	3	−0.15
3	5	−0.19

These results can be conveniently interpreted. For example, at time $t = 5$ the increase of x_2 from 0 to 1 leads to an average change in the Survival probability of −0.19, where Survival Probability is in the range 0-1. Additionally, as mentioned in the previous sections, all estimates of the model will first be yielded as posterior predictive distributions. The empirical mean and percentile function of this distribution yield the point estimates described above and their respective credible intervals.

Demonstration

The following is a brief demonstration on how to use BART-Survival. The *rossi* survival dataset from the lifelines library is used for this example.

```
from lifelines.datasets import load_rossi
from bart_survival import surv_bart as sb
import numpy as np
```

```
#####
# Load rossi dataset from lifelines
rossi = load_rossi()
names = rossi.columns.to_numpy()
rossi = rossi.to_numpy()
```

252 After loading the libraries and data, the first step is to generate the **TAD** and **PAD** datasets.
 253 In this step, the time (originally in days) is downscaled by a factor of 7, setting time to be
 254 measured in weeks.

```
#####
# Transform data into 'augmented' dataset
# Requires creation of the training dataset and a predictive dataset for inference
# TAD
trn = sb.get_surv_pre_train(
    y_time=rossi[:,0],
    y_status=rossi[:,1],
    x = rossi[:,2:],
    time_scale=7
)

# PAD
post_test = sb.get_posterior_test(
    y_time=rossi[:,0],
    y_status=rossi[:,1],
    x = rossi[:,2:],
    time_scale=7
)
```

255 The next step is to initialize the model, which involves setting several parameter values. The key
 256 considerations when initializing the models is number of trees and the split rules. The number of
 257 trees controls how many regression trees will be used. Typically 50 trees is a good default, but
 258 it can be adjusted to assist in model performance. Split rules is a specific PyMC-BART parameter
 259 and is used to designate the how the regression trees are constructed. The one requirement
 260 of the split rules is that the time covariate has to be set as a `pmb.ContinuousSplitRule()`.
 261 Otherwise, generally continuous variables can assigned `pmb.ContinuousSplitRule()` and
 262 categorical variables assigned `pmb.OneHotSplitRule()`. It is recommended to review the
 263 PyMC-BART literature for more information regarding parameterization of the models.

```
#####
# Instantiate the BART models
# model_dict defines specific model parameters
model_dict = {"trees": 50,
              "split_rules": [
                  "pmb.ContinuousSplitRule()", # time
                  "pmb.OneHotSplitRule()", # fin
                  "pmb.ContinuousSplitRule()", # age
                  "pmb.OneHotSplitRule()", # race
                  "pmb.OneHotSplitRule()", # wexp
                  "pmb.OneHotSplitRule()", # mar
                  "pmb.OneHotSplitRule()", # paro
                  "pmb.ContinuousSplitRule()", # prio
              ]
            }
# sampler_dict defines specific sampling parameters
sampler_dict = {
    "draws": 200,
    "tune": 200,
    "cores": 8,
    "chains": 8,
    "compute_convergence_checks": False
}
```

```
BSM = sb.BartSurvModel(
    model_config=model_dict,
    sampler_config=sampler_dict
)
```

264 The model can then be trained with the **TAD** input and predicted p_{ij} values yielded with the
265 **PAD** input.

```
#####
```

```
# Fit Model with TAD
```

```
BSM.fit(
    y = trn["y"],
    X = trn["x"],
    weights=trn["w"],
    coords = trn["coord"],
    random_seed=5
)
```

```
# Get posterior predictive for evaluation using the PAD
```

```
post1 = BSM.sample_posterior_predictive(
    X_pred=post_test["post_x"],
    coords=post_test["coords"]
)
```

266 Finally the Survival probability can derived from the p_{ij} estimates.

```
# Convert to SV probability.
```

```
sv_prob = sb.get_sv_prob(post1)
print(sv_prob["sv"].shape)
# (1600, 432, 8)
```

267 The sv_prob object above is a dictionary containing numpy arrays of both the p_{ij} and s_{ij}
268 estimates, labeled "prob" and "sv" respectively. The p, S arrays are three dimensional with the
269 dimensions of the arrays being: - axis 0 = draws of the posterior predictive distribution: 1600 -
270 axis 1 = observations i : 432 - axis 2 = times j : 8

271 These arrays can be easily reduced down to point estimates and credible intervals using basic
272 numpy methods. For example to get the estimate of the mean over all observations, first get
273 the mean over the observations (axis 1) followed by the mean over the posterior draws (axis
274 0). The results being the estimated mean Survival over the 8 time intervals.

275 Similarly the 0.05-0.95 credible interval for the estimated mean Survival can be returned as
276 the quantile evaluations of the same mean-over-axis-1 array. This yields a (2,8) array with the
277 lower and upper bounds (rows) of the credible interval defined for each time point (columns).

```
# get the mean value across observations for each time within each draw of the posterior
ave_obs = sv_prob["sv"].mean(axis=1)
print(ave_obs.shape)
# (1600, 8)
```

```
# get the average across the posterior draws
```

```
ave_obs_draws = ave_obs.mean(0)
print(ave_obs_draws)
# [0.97259158 0.93862734 0.9015028 0.86572104 0.83166388 0.79712983 0.76321937 0.736341
```

```
# get the .05 and .95 percentiles of the mean across posterior draws
```

```
ci_obs_draws = np.quantile(ave_obs, [0.05, 0.95], axis=0)
print(ci_obs_draws)
# lower bound
```

```
#[[0.96432971 0.92615903 0.88601098 0.84845099 0.81381525 0.77771897 0.7426666 0.714382
# upper bound
# [0.98066217 0.94989412 0.91620586 0.8829086 0.84948681 0.81667259 0.78434529 0.757959
```

278 Examples of generation of marginal effect estimates can be found in the example notebooks
279 provided in the repository documentation.

280 Conclusion

281 BART-Survival provides the computational methods required for completing non-parametric
282 discrete-time Survival analysis. This approach can have several advantages over alternative
283 Survival methods. These advantages include capabilities to incorporate non-linear and interac-
284 tion effects into the model, naturally ability to regularize the model (which reduces the risk of
285 over-fitting) and of being robust to issues of multi-collinearity. The BART-Survival approach
286 is especially useful when the assumptions of alternative Survival methods are violated.

287 Our BART-Survival algorithm has been tested in a rigorous simulation study, with additional
288 applications to real-world data. While the manuscript for this work is currently under devel-
289 opment, the results indicate similar performance as the the R-based BART Survival method
290 across settings of varied complexity. Both methods demonstrate the previously describe ad-
291 vantages over other survival approaches (such as Cox Proportional Hazard Models) when the
292 relationships within the data becomes more complex or assumptions of the these other models
293 are violated. An example comparing the R-based method and our BART-Survival algorithm is
294 provided [here](#).

295 Our library provides a convenient API for completing discrete-time Survival analysis, along with
296 the functionality to customize the methodology as needed. The associated API documentation
297 can be found [here](#), along with the associated github repository [BART-Survival](#).

298 Acknowledgements

299 We thank Oscar Rincón-Guevara for helpful suggestions and review. We also thank Tegan
300 Boehmer, Sachin Agnihotri, and Matt Ritchey for supporting the project throughout its
301 development.

302 References

- 303 Abril-Pla, O., Andreani, V., Carroll, C., Dong, L., Fonnesbeck, C. J., Kochurov, M., Kumar,
304 R., Lao, J., Luhmann, C. C., Martin, O. A., Osthege, M., Vieira, R., Wiecki, T., & Zinkov,
305 R. (2023). PyMC: A modern, and comprehensive probabilistic programming framework in
306 Python. *PeerJ Computer Science*, 9, e1516. <https://doi.org/10.7717/peerj-cs.1516>
- 307 Altman, D. G., & Bland, J. M. (1998). Statistics Notes: Time to event (survival) data. *BMJ*,
308 317(7156), 468–469. <https://doi.org/10.1136/bmj.317.7156.468>
- 309 Bradburn, M. J., Clark, T. G., Love, S. B., & Altman, D. G. (2003). Survival Analysis Part II:
310 Multivariate data analysis – an introduction to concepts and methods. *British Journal of*
311 *Cancer*, 89(3), 431–436. <https://doi.org/10.1038/sj.bjc.6601119>
- 312 Chipman, H. A., George, E. I., & McCulloch, R. E. (2010). BART: Bayesian additive regression
313 trees. *The Annals of Applied Statistics*, 4(1). <https://doi.org/10.1214/09-AOAS285>
- 314 Cox, D. R. (1972). Regression Models and Life-Tables. *Journal of the Royal Statistical Society*
315 *Series B: Statistical Methodology*, 34(2), 187–202. <https://doi.org/10.1111/j.2517-6161.1972.tb00899.x>
- 316

- 317 Harrell, F. E. (2015). *Regression Modeling Strategies: With Applications to Linear Models,*
318 *Logistic and Ordinal Regression, and Survival Analysis*. Springer International Publishing.
319 <https://doi.org/10.1007/978-3-319-19425-7>
- 320 Ishwaran, H., Kogalur, U. B., Blackstone, E. H., & Lauer, M. S. (2008). Random survival
321 forests. *The Annals of Applied Statistics*, 2(3). <https://doi.org/10.1214/08-AOAS169>
- 322 Joffe, E., Coombes, K. R., Qiu, Y. H., Yoo, S. Y., Zhang, N., Bernstam, E. V., & Kornblau, S.
323 M. (2013). Survival Prediction In High Dimensional Datasets – Comparative Evaluation
324 Of Lasso Regularization and Random Survival Forests. *Blood*, 122(21), 1728–1728.
325 <https://doi.org/10.1182/blood.V122.21.1728.1728>
- 326 Quiroga, M., Garay, P. G., Alonso, J. M., Loyola, J. M., & Martin, O. A. (2023). *Bayesian*
327 *additive regression trees for probabilistic programming* (No. arXiv:2206.03619). arXiv.
328 <https://doi.org/10.48550/arXiv.2206.03619>
- 329 Sparapani, R. A., Logan, B. R., McCulloch, R. E., & Laud, P. W. (2016). Nonparametric
330 survival analysis using Bayesian Additive Regression Trees (BART). *Statistics in Medicine*,
331 35(16), 2741–2753. <https://doi.org/10.1002/sim.6893>
- 332 Sparapani, R., Spanbauer, C., & McCulloch, R. (2021). Nonparametric Machine Learning and
333 Efficient Computation with Bayesian Additive Regression Trees: The **BART** R Package.
334 *Journal of Statistical Software*, 97(1). <https://doi.org/10.18637/jss.v097.i01>

DRAFT